

Computing Derivatives

Section I. Truss Optimization

For this assignment, objective function derivatives and constraint derivatives were computed and passed to MATLAB's 'fmincon' function in order to optimize the classic Ten Bar Truss problem. Three different derivative computing methods were implemented namely forward differencing, central differencing, and complex step. Findings and observations from this assignment are given in this report.

Part a. Is this a problem that could benefit from scaling?

	Forward Difference		Central Difference		Complex Step	
No Scaling	nfun = 991	etime = 0.40	nfun = 2773	etime = 0.61	nfun = 1849	etime = 0.87
Scaled 1e-3	nfun = 331	etime = 0.14	nfun = 1513	etime = 0.43	nfun = 859	etime = 0.40

As can be seen from the table above, this is a problem that benefits significantly from scaling. As can be seen in the case of forward difference, the number of function calls and execution time is **three times less** than it is for the unscaled problem. Similarly central difference and the complex step methods have significantly fewer function calls and reduced execution time after scaling is applied. The benefits of scaling were expected due to the large difference in magnitude between the values for weight and the values for stress in the truss problem. By dividing the stresses by 1000 and bringing them down closer to the order of magnitude of the weights, the optimization operates much more efficiently. Note: A step size of 1e-4 was used for derivative calculations in all three methods.

Part b. Listings of MATLAB code

Listed below are the sections of MATLAB code where derivatives were computed for each of the three methods.

Forward Difference:

```
function [f, gradf] = obj(x)
    [f, ~, ~] = objcon(x);
    % compute forward difference objective derivatives
    delta_x = 1e-4; % set the step size
    gradf = zeros(length(x),1); % initialize gradient vector
    % perturb each element of x and compute the partial derivatives
    for i = 1:length(x)
        x_p = x; % reset x_p
        x_p(i) = x_p(i) + delta_x; % perturb x_p
        [f_p, ~, ~] = objcon(x_p); % compute objective at perturbed x
        gradf(i) = (f_p - f)/delta_x; % compute forward diff derivative
    end
    % myobjgrad = gradf
end
```

```

function [c, ceq, DC, DCeq] = con(x)
    [~, c, ceq] = objcon(x);
    % compute forward difference constraint derivatives
    delta_x = 1e-4; % set the step size
    DC = zeros(length(x)); % initialize gradient vector
    % perturb each element of x and compute the partial derivatives
    for i = 1:length(x)
        x_p = x; % reset x_p
        x_p(i) = x_p(i) + delta_x; % perturb x_p
        [~, c_p, ~] = objcon(x_p); % compute constraint at perturbed x
        DC(i,:) = (c_p - c)/delta_x; % compute forward diff derivative
        DCeq = [];
    end
    % mycongrad = DC
end

```

For the forward difference method, I simply integrated the calculation of derivatives into the obj() and con() functions that we pass to fmincon(). This approach is nice because we only have to compute the unperturbed objective value once and we can use it for both the objective function value that we pass to fmincon, as well as for the computation of the derivative. The only difficulty encountered with this method was that the constraint derivatives initially would not pass the fmincon 'CheckGradients' routine—regardless of how the step size was chosen. However after scaling was applied, and the step size for the constraint derivatives was set to 1e-6, the gradients passed without any issue. The forward difference method computes the derivative by evaluating the function at a small step forward from the current point, and then computing the slope between the current point and the forward point. This method works well enough if the step size is chosen appropriately. If the step size is too large, truncation error will creep in, and if the step size is too small, roundoff error due to subtractive cancellation will cause problems. For this method, the expected truncation error is proportional to the step size—i.e., **expected error is on the order of delta_x**.

Central Difference:

```

function [f, gradf] = obj(x)
    [f, ~, ~] = objcon(x);
    % compute central difference objective derivatives
    delta_x = 1e-4; % set the step size
    gradf = zeros(length(x),1); % initialize gradient vector
    % perturb each element of x and compute the partial derivatives
    for i = 1:length(x)
        x_p = x; % reset x_p
        x_m = x;
        x_p(i) = x_p(i) + delta_x; % perturb x_p
        x_m(i) = x_m(i) - delta_x;
        [f_p, ~, ~] = objcon(x_p); % compute objective at perturbed x
        [f_m, ~, ~] = objcon(x_m); % compute objective at perturbed x
        gradf(i) = (f_p - f_m)/(2*delta_x); % compute central diff derivative
    end
end

function [c, ceq, DC, DCeq] = con(x)
    [~, c, ceq] = objcon(x);
    % compute central difference constraint derivatives
    delta_x = 1e-4; % set the step size
    DC = zeros(length(x)); % initialize gradient vector
    % perturb each element of x and compute the partial derivatives

```

```

for i = 1:length(x)
    x_p = x; % reset x_p
    x_m = x;
    x_p(i) = x_p(i) + delta_x; % perturb x_p
    x_m(i) = x_m(i) - delta_x;
    [~, c_p, ~] = objcon(x_p); % compute constraint at perturbed x
    [~, c_m, ~] = objcon(x_m); % compute constraint at perturbed x
    DC(i,:) = (c_p - c_m)/(2*delta_x); % compute forward diff derivative
    DCeq = [];
end
end

```

Derivatives for the central difference method were incorporated into the code the same way as they were for forward difference. The only major difference here is that the objective function must be called within the obj() and con() functions three times instead of twice to allow for perturbing the function forward a step and backwards a step. The central difference method computes the derivative by evaluating the function at a small step backwards **and** forwards of the current point, and then computing the slope between the backward point and the forward point. No problems were encountered for this method and 'CheckGradients' passed without any problems. **Central differencing is more accurate than forward differencing and the expected error is on the order of δx^2 .** The main drawback to this method is that it is more computationally heavy since additional function evaluations are required.

Complex Step:

```

function [f, gradf] = obj(x)
    [f, ~, ~] = objcon(x);

    % compute the partial objective derivatives via complex step

    delta_x = 1e-4; % set the step size
    gradf = zeros(length(x),1); % initialize gradient vector
    for it = 1:length(x)
        x_p = x;
        x_p(it) = x_p(it) + delta_x*i;
        [f_p, ~, ~] = objcon(x_p);
        gradf(it) = imag(f_p)/delta_x;
    end
end
function [c, ceq, DC, DCeq] = con(x)
    [~, c, ceq] = objcon(x);

    % compute the partial constraint derivatives via complex step

    delta_x = 1e-4; % set the step size
    DC = zeros(length(c),length(x)); % initialize gradient vector
    for it = 1:length(x)
        x_p = x;
        x_p(it) = x_p(it) + delta_x*i;
        [~, c_p, ~] = objcon(x_p);
        DC(:, it) = imag(c_p)/delta_x;
    end
    DC = DC';
    DCeq = [];
end

```

Once again, derivatives were incorporated into the code the same way as the previous two methods. The only challenge encountered was changing how the constraint functions were defined within the `objcon()` function. Since the absolute value function isn't defined for complex numbers, constraints were split into 20 separate constraints instead of 10. For this method, derivatives are computed by perturbing the function by a complex `delta_x` as is shown in Eq. 4.18 from the notes. This method is really slick and cleverly avoids the ill effects of subtractive cancellation. Essentially you can choose a `delta_x` for this method to be as small as you like, and still compute exceptionally accurate derivatives. Similar to the central difference method, expected error is on the order of delta_x^2 but now we can make `delta_x` to be very very small.

Part c. Error Evaluation

Since weight scales linearly in this problem, accuracy of derivatives isn't dependent on step size and so all three methods produced equally accurate derivatives for weight. The constraint derivatives related to stresses however are sensitive to step size and so accuracy between the three methods could be evaluated. Since the complex step method allows us to pick a very very small step size (`delta_x`), for this evaluation, a step size of $1\text{e-}20$ was selected for the complex step, and the constraint derivatives for this method were treated as the 'true' derivatives for comparison. For all three methods, the constraint derivatives from the starting point were collected and then compared to the 'true' derivative generated by the complex step approach. For the forward difference and central difference methods a step size of $1\text{e-}4$ was chosen. This step size was chosen because it gave sufficiently small error in the derivatives and could be used to optimize the truss problem without issue.

To evaluate the errors, I simply took the norm of the difference between the complex step constraint derivatives and the forward/central difference constraint derivatives. See code snippet below.

```
error_fd = norm(complexDC.DC - fdDC.DC)
error_cd = norm(complexDC.DC - cdDC.DC)
error_complex = norm(complexDC.DC - complexDC.DC)
```

Result:

```
error_fd =
```

```
2.0060e-04
```

```
error_cd =
```

```
3.9676e-09
```

```
error_complex =
```

```
0
```

As expected and as can be seen, the forward difference method has the most error and is on the order of `delta_x`. The central difference method has significantly less error even though the same step size was

used and its error is approximately on the order of Δx^2 . Error in the complex step here is assumed zero. Each of the three methods has its merits. Forward difference is fast (computationally) and calculation of derivatives is the most straightforward. Central differencing takes approximately twice as much computation as forward differencing but for the same step size in this case has error that is five orders of magnitude smaller. That's pretty significant! The complex step method is the most accurate but is not a 'blackbox' method like the other two are and requires that the objective function can handle complex values. Additionally, it costs more computationally to perform operations on complex variables. Depending on the nature of the function, and the constraints on computation, there are trade offs of speed vs accuracy.

Part d. Table

	Forward Difference	Central Difference	Complex Step
Number of Function Calls	397	1639	859
Time of Execution	0.20	0.35	0.37
Stopping Criterion	* below	* below	* below
Final Objective Value	1.5932e+03	1.5932e+03	1.5932e+03

Note: Data in the above table is *after* appropriate scaling was applied and time of execution represents the average of three trial runs.

* The relative first-order optimality measure, 1.295203e-08, is less than options.OptimalityTolerance = 1.000000e-06, and the relative maximum constraint violation, 0.000000e+00, is less than options.ConstraintTolerance = 1.000000e-06.

Discussion

The biggest difference that is observed in the data above is the number of function calls, and the time of execution. As expected, forward difference requires the fewest function evaluations and as a result, requires the least amount of time to execute. Central difference requires the most function evaluations and the time of execution is increased. Complex step takes more evaluations than forward difference, and less than complex step but takes the most time to execute. This increased execution time results from the extra computational resources that are required to evaluate complex variables.

Section II. Automatic Differentiation for the Spring Problem

Part a. Explanation

The main mathematical mechanism of the Automatic Differentiation method (AD) is the chain rule for differentiation. In this method, function values are computed one component at a time such that each line in the program is composed of the values computed on previous lines. In other words, function derivatives are automatically computed as part of computing the actual function value. Then using object oriented programming, and function overloading, the function values and function derivatives can be computed simultaneously.

Part b. MATLAB code applying valder class to the Spring Optimization Problem

```
clc
clear
close all

x = [0.015, 0.5, 10.0, 1.5];

[function_values, function_jacobians, constraint_values, constraint_jacobians] = objcon(x);

function_values
function_jacobians

constraint_values
constraint_jacobians

% -----Objective and Non-linear Constraints-----
function [function_vals, function_jacobians, constraint_vals, constraint_jacobians] = objcon(x)

% set objective/constraints here

% make design variables valder objects (things we'll adjust to find optimum)
d = valder(x(1),[1,0,0,0]); % wire dia (in)
D = valder(x(2),[0,1,0,0]); % coil dia (in)
n = valder(x(3),[0,0,1,0]); % num coils
hf = valder(x(4),[0,0,0,1]); % free height (no load) (in)

% other analysis variables (constants that the optimization won't touch)
h0 = 1.0; % preloaded height (in)
delta0 = 0.4; % deflection (in)
hdef = h0 - delta0; % deflected spring height (in)
G = 12e6;
Q = 150e3;
w = 0.18;
Se = 45e3;
Sf = 1.5;

% delta_x = 0.4; % not sure if this is right??
delta_x = (hf - h0); % maybe this instead dc = zeros(9,1);???

% analysis functions
k = G*d^4/(8*D^3*n);
F = k*delta_x;
```

```

K = ((4*D-d)/(4*(D-d)))+0.62*(d/D);
% Tau = (8*F*D/pi*d^3)*K;
hs = n*d;
F_min = k*(hf - h0);
% F_max = F_min + delta0*k;
F_max = k*(hf - (h0 - delta0));
F_hs = k*(hf - hs);
Tau_min = 8*F_min*D*K/(pi*(d^3));
Tau_max = 8*F_max*D*K/(pi*(d^3));
Tau_m = (Tau_max + Tau_min)/2;
Tau_a = (Tau_max - Tau_min)/2;
Tau_hs = 8*F_hs*D*K/(pi*(d^3));
Sy = 0.44*(Q/d^w);

function_vals = [k.val, F.val, K.val, hs.val, F_min.val, F_max.val, ...
                F_hs.val, Tau_min.val, Tau_max.val, Tau_m.val, Tau_a.val, ...
                Tau_hs.val, Sy.val];
function_jacobians = [k.der; F.der; K.der; hs.der; F_min.der; F_max.der; ...
                    F_hs.der; Tau_min.der; Tau_max.der; Tau_m.der; Tau_a.der; ...
                    Tau_hs.der; Sy.der];

% objective function (what we're trying to optimize)
% f = -F; % maximize Force

% inequality constraints (c<=0)
c1 = Tau_hs - Sy;
c2 = Tau_a - Se/Sf;
c3 = Tau_a + Tau_m - Sy/Sf;
c4 = (D/d) - 16;
c5 = -(D/d) + 4;
c6 = d - 0.2;
c7 = -d + 0.01;
c8 = D + d - 0.75;
c9 = -hdef + hs + 0.05;

constraint_vals = [c1.val, c2.val, c3.val, c4.val, c5.val,...
                  c6.val, c7.val, c8.val, c9.val];

constraint_jacobians = [c1.der; c2.der; c3.der; c4.der; c5.der;...
                      c6.der; c7.der; c8.der; c9.der];

end

```

Result:

```

function_values =

    1.0e+05 *

    0.000000607500000
    0.000000303750000
    0.000010417958763
    0.000001500000000
    0.000000303750000
    0.000000546750000
    0.000000820125000
    0.119381013650947

```

```
0.214885824571704
0.167133419111326
0.047752405460379
0.322328736857557
1.405553914176020
```

```
function_jacobians =
```

```
1.0e+06 *
```

```
0.000016200000000 -0.000000364500000 -0.000000006075000 0
0.000008100000000 -0.000000182250000 -0.000000003037500 0.000000060750000
0.000002834218302 -0.000000085026549 0 0
0.000010000000000 0 0.000000015000000 0
0.000008100000000 -0.000000182250000 -0.000000003037500 0.000000060750000
0.000014580000000 -0.000000328050000 -0.000000005467500 0.000000060750000
0.000021262500000 -0.000000492075000 -0.000000009112500 0.000000060750000
0.828351173720028 -0.048726737941790 -0.001193810136509 0.023876202730189
1.491032112696050 -0.087708128295222 -0.002148858245717 0.023876202730189
1.159691643208040 -0.068217433118506 -0.001671334191113 0.023876202730189
0.331340469488011 -0.019490695176716 -0.000477524054604 0
1.997786141742184 -0.131562192442834 -0.003581430409528 0.023876202730189
-1.686664697011224 0 0 0
```

```
constraint_values =
```

```
1.0e+05 *
```

```
-1.083225177318463
-0.252247594539621
-0.722150118212309
0.0001733333333333
-0.0002933333333333
-0.000001850000000
-0.000000050000000
-0.000002350000000
-0.000004000000000
```

```
constraint_jacobians =
```

```
1.0e+06 *
```

```
3.684450838753409 -0.131562192442834 -0.003581430409528 0.023876202730189
0.331340469488011 -0.019490695176716 -0.000477524054604 0
2.615475244036867 -0.087708128295222 -0.002148858245717 0.023876202730189
-0.002222222222222 0.0000666666666667 0 0
0.002222222222222 -0.0000666666666667 0 0
0.000001000000000 0 0 0
-0.000001000000000 0 0 0
0.000001000000000 0.000001000000000 0 0
0.000010000000000 0 0.000000015000000 0
```


Part c. How AD differs from other numerical methods

One way that AD differs from other methods is that it requires function variables to be 'objects' (in the sense of object oriented programming) that can be overloaded to handle the computation of both values and derivatives simultaneously. AD is also different than traditional numerical derivative computation in that there is no approximation error. This arises from the fact that we're not perturbing the function to get approximate derivatives but instead we are composing function values by computing its derivatives along the way. Finally similar to the complex step, AD is not a blackbox method and requires that the analysis software be modified and executed in an environment that supports AD (such as MATLAB using the valder class). Pros of AD include simultaneous values and derivatives, and no approximation error. The primary con is that it's not supported in all environments, and it requires more computational resources.